



An Introduction to SPARQL



The world's most performant knowledge graph and semantic reasoning engine.

Objectives



Setting up RDFox

- ☐ License and executable
- ☐ IDE
- ☐ REST endpoint and UI

Loading data into RDFox

- ☐ Create datastore
- ☐ Import data

Exploring and querying with SPARQL

- ☐ Basics of SPARQL
- ☐ Useful queries to explore
- ☐ Aggregates
- ☐ Negation
- ☐ Filters
- ☐ Binds
- ☐ Optionals

Please feel free to ask questions at any time!
Exercises are scattered throughout.
Links are provided for reference.



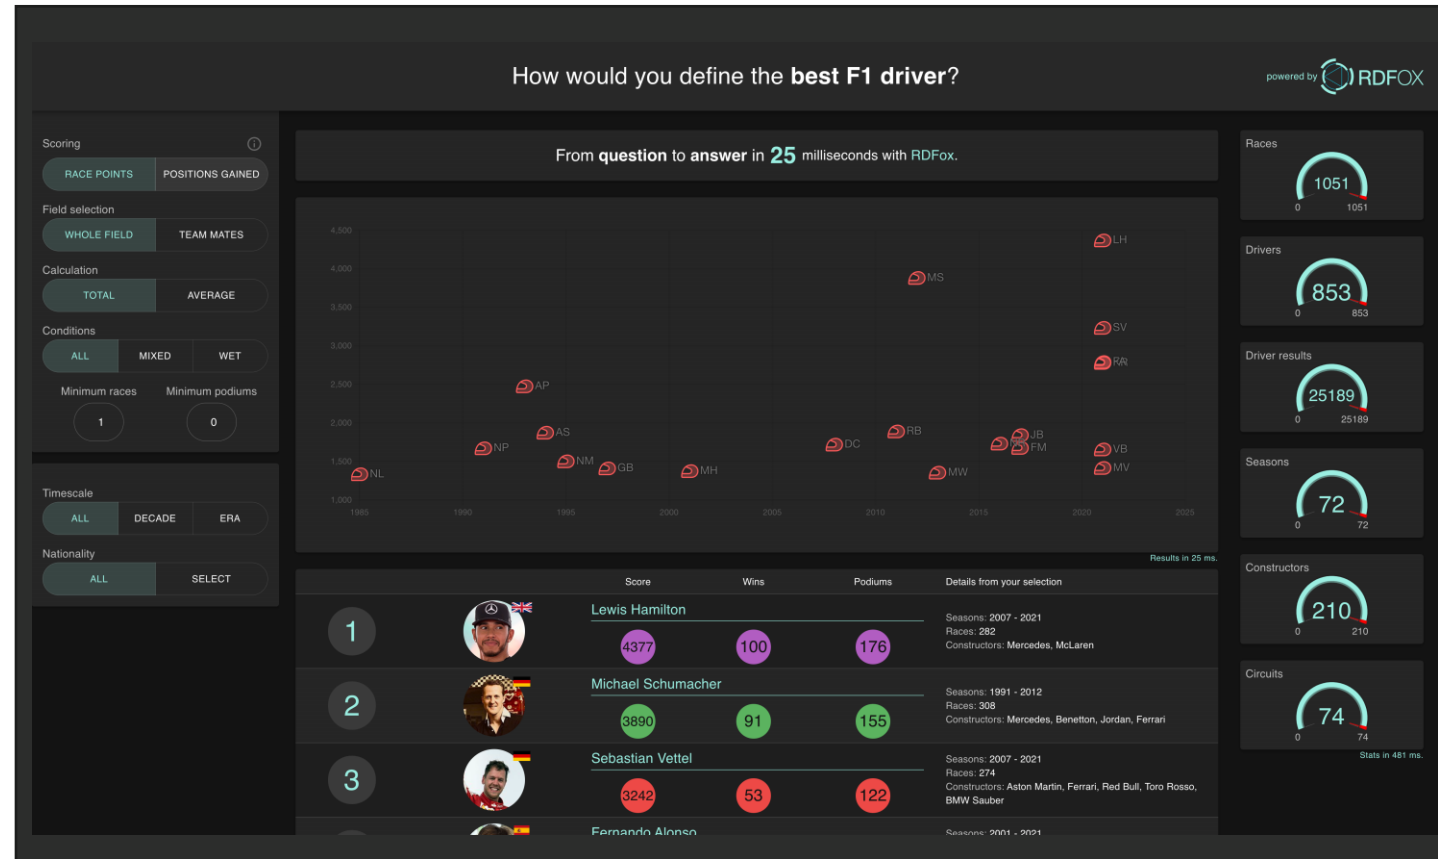
<https://docs.oxfordsemantic.tech/>



What are you making?

Who is the Greatest Formula One Driver of All Time?

Controls
and filters
for the
scoring
system.



Statistics
about the
data used
to form the
results.

This matches
the filters.

Try it for yourself!

<http://f1.rdfox.tech>

Requirements



A. Get an RDFox License

<https://www.oxfordsemantic.tech/tryrdfoxforfree>

B. Download RDFox (& unzip)

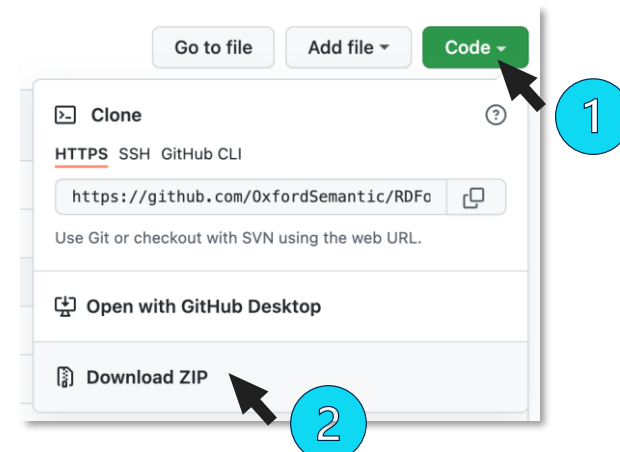
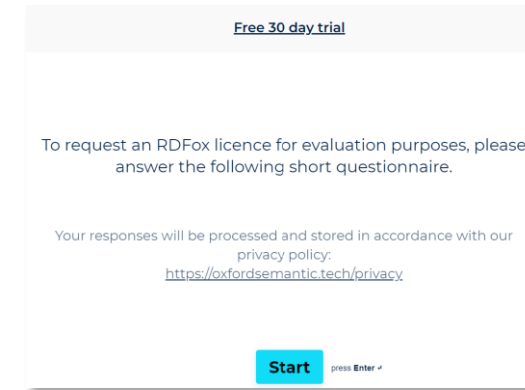
<https://www.oxfordsemantic.tech/downloads>


C. Download the class materials from Github:

<https://github.com/OxfordSemantic/RDFoxWorkshop-SPARQL>

OPTIONAL Get your IDE of choice ready (VS Code etc.)

<https://code.visualstudio.com/>





You should have...

examples	✓	30/03/2022 11:23	File folder
include	✓	30/03/2022 11:23	File folder
lib	✓	30/03/2022 11:23	File folder
RDFox.lic	✓	30/03/2022 11:23	Text Document
RDFox	✓	30/03/2022 11:23	Application
readme	✓	30/03/2022 11:23	Text Document
version	✓	30/03/2022 11:23	Text Document



Make sure to put your license
INSIDE the **RDFox folder**.

Open the class PDF so you can follow along!
You will need to follow the links provided.

Queries	✓	26/04/2022 16:42	File folder	
.DS_Store	⊖	26/04/2022 16:42	DS_STORE File	7 KB
.gitignore	✓	26/04/2022 16:42	Git Ignore Source ...	1 KB
LICENSE	✓	26/04/2022 16:48	File	2 KB
RDFoxWorkshop-SPARQL	↻	05/04/2022 18:30	Chrom	
README	✓	26/04/2022 16:42	Marko	
upTo2020.ttl	✓	26/04/2022 16:42	Comp	



You should have...



☐ The class files



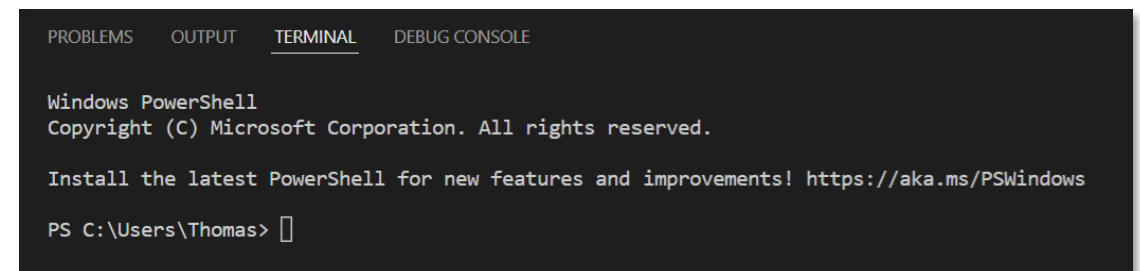
☐ RDFox & License



☐ This presentation



☐ Your IDE Open



Setting Up RDFox

Setting up RDFox

- ☐ License and executable
- ☐ IDE
- ☐ REST endpoint and UI

Loading data into RDFox

- ☐ Create datastore
- ☐ Import data

Exploring and querying with SPARQL

- ☐ Basics of SPARQL
- ☐ Useful queries to explore
- ☐ Aggregates
- ☐ Negation
- ☐ Filters



Setting up RDFox

- Open your favourite IDE, and open the RDFoxClass folder
- Open a terminal in the IDE
- In the terminal, run the following
Mac: `<Path/to>/RDFox sandbox`
Windows: `<Path/to>/RDFox.exe sandbox`
- The RDFox server is now running, but it's empty...

If your file path contains spaces put "quotation marks" around it.

You may need to use back slashes \ or double up the slashes // depending on your setup.



```
Source code for RDFox v1.0 Copyright 2013 Oxford University Innovation Limited and subsequent improvements Copyright 2017-2021 by Oxford Semantic Technologies Limited.
```

```
This copy of RDFox is licensed for Developer use to Tom Vout (tom.vout@oxfordsemantic.tech) of OST until 07-Jun-2022 16:01:44
```

```
This system is equipped with 16.9 GB of RAM, and RDFox is configured to use at most 15.2 GB (89.9% of the total).
```

```
Currently, 2.8 GB (18.4% of the amount allocated to RDFox) appear to be available on the system.
```

```
Since RDFox is a RAM-based system, its performance can suffer when other running processes use a lot of memory.
```

```
A new server connection was opened as role 'guest' and stored with name 'sc1'.  
> 
```



<https://docs.oxfordsemantic.tech/command-line-reference.html#starting-the-rdfox-process>



The REST Endpoint and UI


- In the terminal, run:
`endpoint start`

```
> endpoint start

WARNING: The RDFox endpoint is running with no transport layer security (TLS). This could allow attackers to
steal information including role passwords.
        See the endpoint.channel variable and related variables in the description of the RDFox endpoint for
        details of how to set up TLS.

The REST endpoint was successfully started at port number/service name 12110 with 7 threads.
> 
```

- The RDFox endpoint is now running, so we can use the web UI
- Open a browser and go to:
localhost:12110/console
- This will show an empty console



<https://docs.oxfordsemantic.tech/command-line-reference.html#endpoint>

Loading Data into RDFox

Setting up RDFox

- ✓ License and executable
- ✓ IDE
- ✓ REST endpoint and UI

Loading data into RDFox

- ❑ Create datastore
- ❑ Import data

Exploring and querying with SPARQL

- ❑ Basics of SPARQL
- ❑ Useful queries to explore
- ❑ Aggregates
- ❑ Negation
- ❑ Filters
- ❑ Binds
- ❑ Optionals



Loading data into RDFOx – Create data store

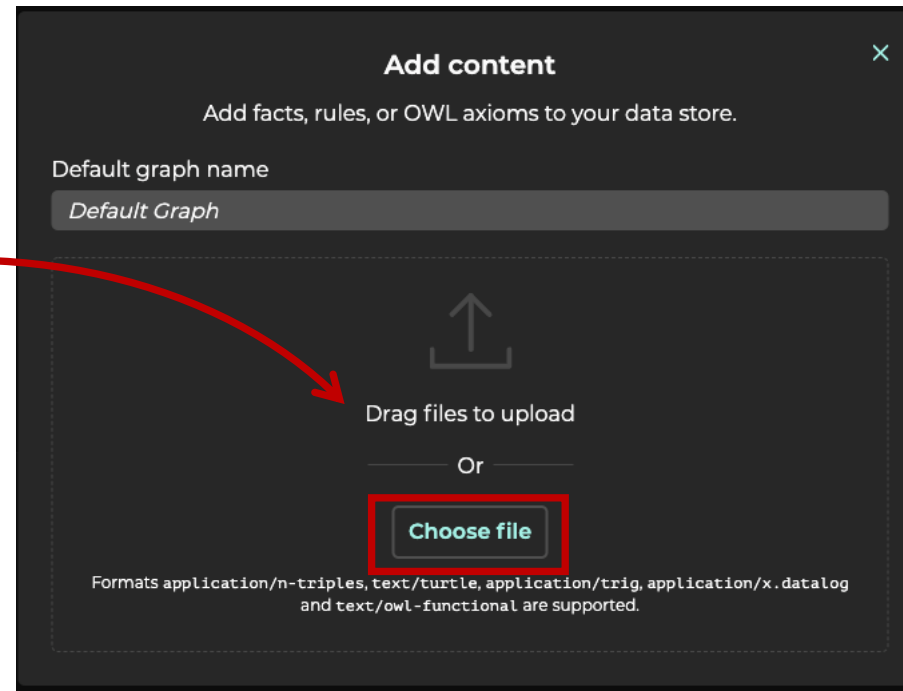

The screenshot illustrates the steps to create a data store in RDFOx. In the main interface, the '+ Create data store' button is highlighted. The 'Create data store' modal dialog shows a 'Name' field containing 'f1' and a 'Create data store' button. The 'Data store has been created' confirmation dialog shows an 'Import content' button.

Make sure to call your data store **f1** so that the links in this guide work.



<https://docs.oxfordsemantic.tech/getting-started.html#getting-started-with-the-web-console>

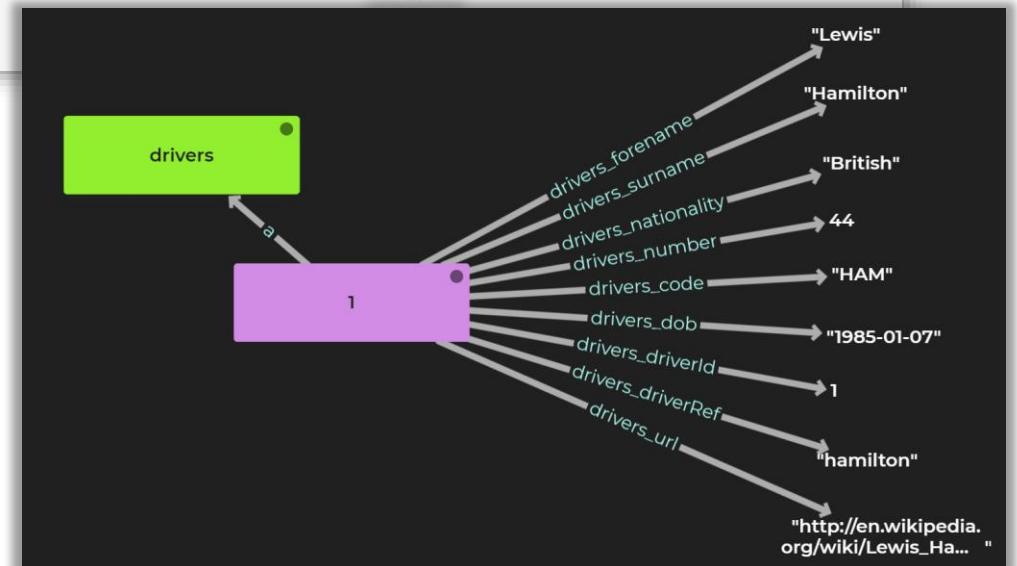
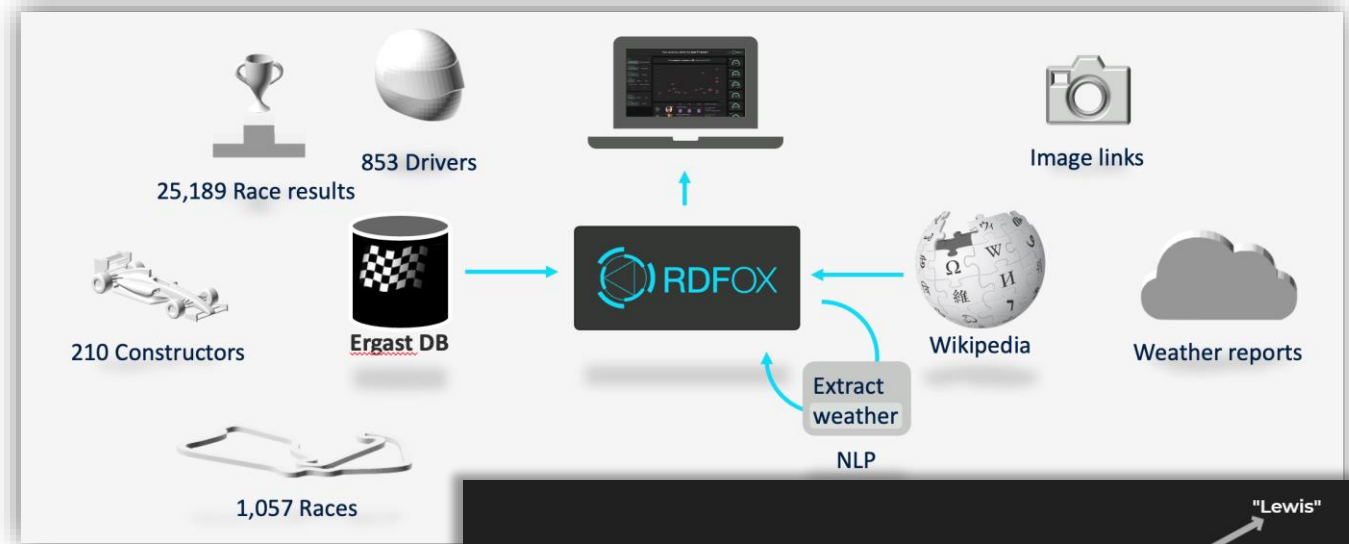
Loading data into RDFox – Load files



The data for today's class



- Race data from the fan-run Ergast database
 - Drivers
 - Races
 - Race results
 - Constructors
- Combined with additional information from Wikidata
 - Driver images
 - Race weather reports



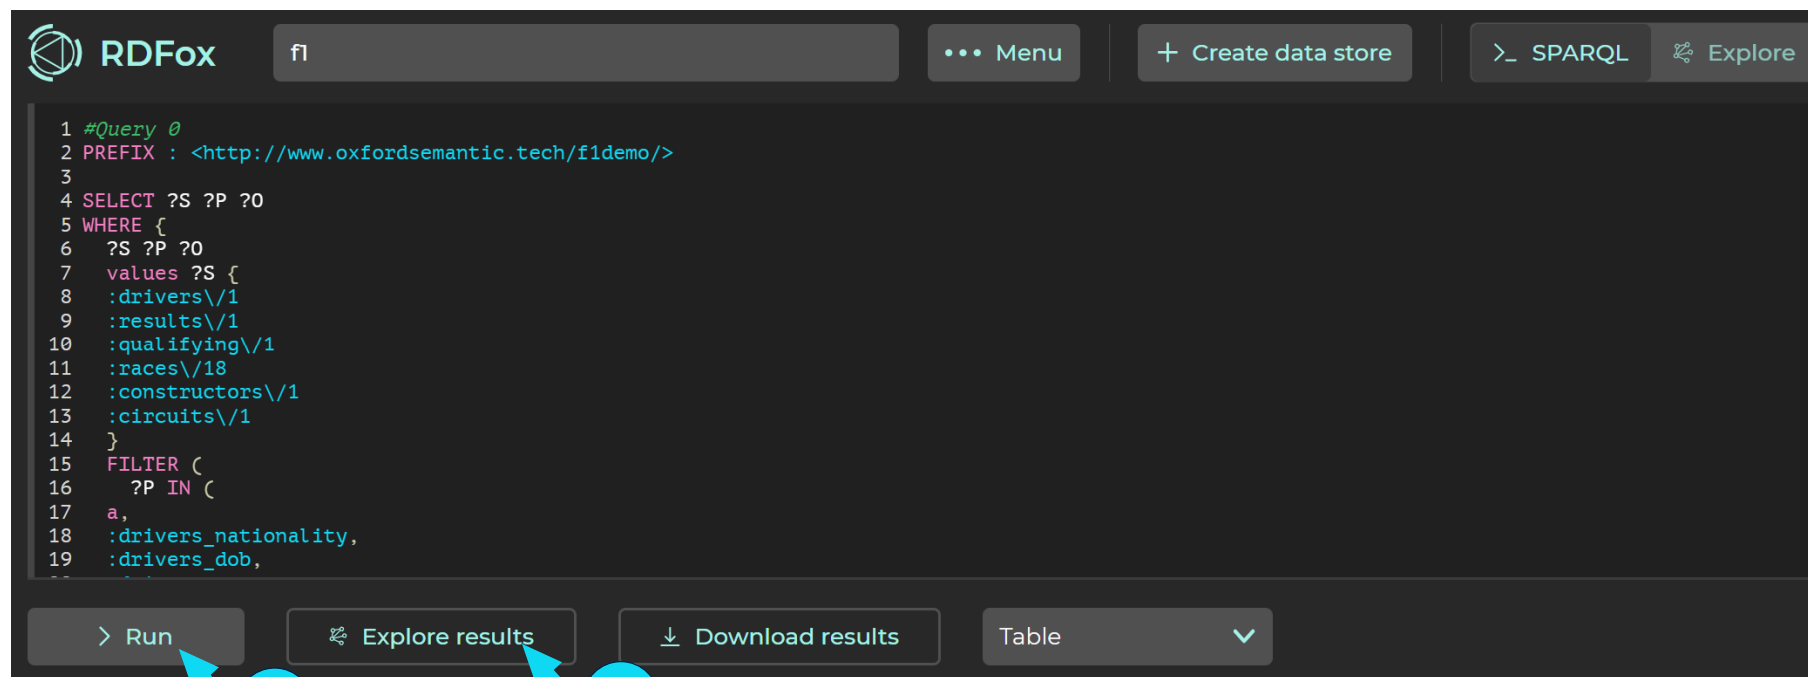
Query 0 – Exploring the data

[Console Query](#)



Click this link to open the query in the console!

Don't worry if it looks complicated, you don't need to understand this query yet!



Click **Run**

and then **Explore**

Exploring and Querying with SPARQL

- ✓ License and executable
- ✓ IDE
- ✓ REST endpoint and UI

Loading data into RDFox

- ✓ Create datastore
- ✓ Import data

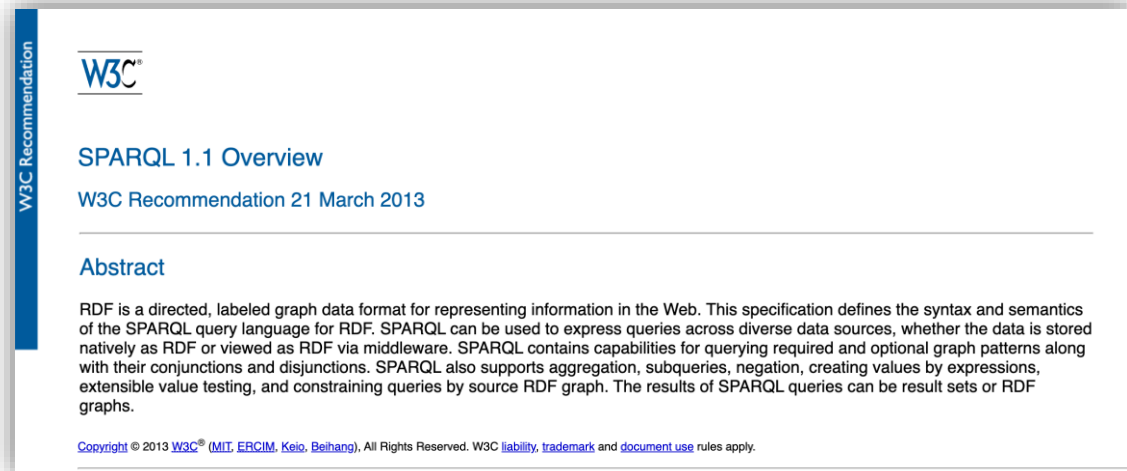
Exploring and querying with SPARQL

- ☐ Basics of SPARQL
- ☐ Useful queries to explore
- ☐ Aggregates
- ☐ Negation
- ☐ Filters
- ☐ Binds
- ☐ Optionals

SPARQL “used to express queries”



<https://www.w3.org/TR/sparql11-overview/>





Query 1 – Basic SELECT

SELECT queries are used to extract results from the dataset and can be modified to match whatever pattern you desire.

This is the simplest **SELECT** query. It returns all the triples in your dataset.

The **SELECT** keyword determines the type of query.

SELECT simply returns values.

```
1 PREFIX : <http://www.oxfordsemantic.tech/f1demo/>
2
3 # This query matches all the triples,
4 # and returns all the variables from the matched results
5 SELECT ?S ?P ?O
6 WHERE {
7   ?S ?P ?O
8 }
```

The output variables **?S ?P** and **?O** tell the **SELECT** query what to return from the variables that are found by the **WHERE** clause.

Here we're asking for everything.

The **WHERE** clause describes what the query is looking for.

In this case we want all triples, so we write the variables **?S ?P** and **?O**, showing we want cases where a subject, predicate, and object exist, with any value in any position.

You can call your variables whatever you like but upper case **?S ?P** and **?O** have some special properties that we will make use of later.



Query 2 – DISTINCT SELECT

Often you'll want to be more specific.

This time, we don't want all the triples to be returned, just their predicates.

Using **SELECT** alone would return duplicates, but we don't want that either.

The **DISTINCT** modifier tells the query to disregard duplicate results.

```
1 PREFIX : <http://www.oxfordsemantic.tech/f1demo/>
2
3 # What kinds of edges are in the data
4 SELECT DISTINCT ?P
5 WHERE {
6   ?S ?P ?O
7 }
```

This time we only ask for **?P** to be returned.

We still have to tell the query that the variable **?P** should come from the predicate of a triple, and that we want to consider all triples.

Therefore, we still need **?S ?P ?O** inside the **WHERE** clause.



Query 3 – Finding Classes

We want to **SELECT** all the **DISTINCT** classes that exist in our data to find the all the types of things we're storing.

In SPARQL, the keyword '**a**' is used universal as the predicate for specifying the type of an entity.

We want to find all instances of triples **WHERE** nodes have the '**a**' property.

```
1 PREFIX : <http://www.oxfordsemantic.tech/f1demo/>
2
3 # What classes of nodes are there in the data?
4 SELECT DISTINCT ?type
5 WHERE {
6   ?S a ?type
7 }
```

We are only interested in the object of these triples – their class.

We have defined the variable **?type** to help us.



Query 4 – Class Properties

This is one of the most important queries we'll cover today.

We **SELECT** a class of interest, and then return all its properties—not their values, just the properties it possesses.

This whole line declares a prefix. This creates a substitute, so every time you see ':', it is equivalent to <http://www.oxfordsemantic.tech/f1demo/>.

We want to find all instances of the class **:drivers**.

```
1 PREFIX : <http://www.oxfordsemantic.tech/f1demo/>
2
3 # What types of edges are there specifically for drivers?
4 SELECT DISTINCT ?p
5 WHERE {
6   ?s a :drivers ;
7   ?p ?o .
8 }
```

Subjects can be shared across multiple lines.

To denote this, all lines except the last must end with a ';'.
The last ends with a '.'.

And to consider all triples that share their subjects, **?s**.

Exercise 1



Run through **queries 1-4** in the **web console**.

Try them out yourself. See what they do and how they do it.



Query 5 – Finding Drivers

Find all drivers along with their forename and surname.

```
1 PREFIX : <http://www.oxfordsemantic.tech/f1demo/>
2
3 # What are the names of my drivers?
4 SELECT DISTINCT ?driver ?forename ?surname
5 WHERE {
6   ?driver a ___ ;
7           :drivers_forename ___ ;
8           :drivers_surname ___ .
9 }
```

We have introduced new variables to the output that must appear in the **WHERE** clause.



Query 6 – A Driver's Properties

To build on what we've got so far, let's find all the associated properties and corresponding values of a specific driver. Say, Lewis Hamilton.

```
1 PREFIX : <http://www.oxfordsemantic.tech/f1demo/>
2
3 # Can you return everything about a specific driver, say Lewis Hamilton?
4 SELECT ?driver ?p ?o
5 WHERE {
6   ?driver a :drivers ;
7           "Lewis" ;
8           "Hamilton" ;
9           .
10 }
```

Literals, or the values attached to properties, can be specified in queries too.

Here we're looking for a specific string so we use the `"` characters to identify it.



Query 7 – Hamilton's Races

Find the total number of races Hamilton has taken part in.

```
1 PREFIX : <http://www.oxfordsemantic.tech/f1demo/>
2
3 # Number of races Lewis Hamilton raced
4 SELECT (COUNT(?race) AS ?raceCount)
5 WHERE {
6
7   ?driver :drivers_forename "Lewis" ;
8           :drivers_surname "Hamilton" .
9
10  ?result :results_driver ___ ;
11         :results_race ___ .
12
13 }
```

The **COUNT** function counts the number of instance of a variable, **?race**, and binds the resulting value as another variable, **?raceCount**.



Query 8 – All Drivers' Races

Find the total number of races each driver individually has taken part in.

```
1 PREFIX : <http://www.oxfordsemantic.tech/f1demo/>
2
3 # Number of races per driver, ordered
4 SELECT ?driver (COUNT(?race) AS ?raceCount)
5 WHERE {
6   ?result :results_driver ___ ;
7           :results_race ?race .
8 }
9 GROUP BY ___
10 ORDER BY DESC(___)
```

The **GROUP BY** clause states how the results should be grouped when returned. Aggregate functions, such as **COUNT**, will now be performed for each group.

The **ORDER BY** clause determines the order of the query results.

The **DESC** modifier tells the **ORDER BY** clause the variable indicated should be ranked in descending order.

Exercise 2



Complete **Queries 5-8**.

HINTS:

It might help you to have **several tabs open** with the previous queries in each one—they will help you fill in the gaps!

You can click on [Incomplete Query](#) at any time to see an editable version of the query for you to play around with.

Sometimes there will be parts missing for you to fill in. You can click on [Answer](#) for the query in full for guidance.



Exercise 2 – Bonus Questions

BONUS Qs

Display the **forename** and **surname** of the driver (**instead of the IRI**) next to their race count.

Answer

Using the **GROUP_CONCAT** aggregate function, find the list of teammates that each driver has ever had.

Answer

HINTS:

Two people are teammates if they competed in the same **race** for the same **constructor**. Notice every result is associated with a constructor and race.

You can put **DISTINCT** inside of an aggregate function to deduplicate.



Query 9 – Drivers Without Podiums

Find all of the drivers who have **never** finished on the podium.

The **FILTER** function contains properties that are used to restrict the results that are returned to just those that match.

```
1 PREFIX : <http://www.oxfordsemantic.tech/f1demo/>
2
3 # Drivers who never got a podium
4 SELECT ____ ____
5 WHERE {
6 # First, get all the drivers with their names.
7   ?driver a :drivers ;
8           :drivers_forename ?forename ;
9           :drivers_surname ?surname .
10
11 # Then make sure that they have never achieved a podium (i.e. positions 3, 2, or 1)
12   FILTER NOT EXISTS {
13     ?result ____ ?driver ;
14             :results_positionOrder ?positionOrder .
15     FILTER(?positionOrder IN(1, 2, 3))
16   }
17 }
```

The **IN** operator checks to see if the variable on the left-hand side appears in the list.

The **NOT EXISTS** operator is our first glimpse of negation. It returns true if its conditions are **not** matched.

In effect, in this case the conditions of the **FILTER** are reversed, so results are limited to those that **do not** match what is stated.



Query 10 – Most Races Without a Podium

Rank all of the drivers who have **never** finished on the podium from most to least races without a podium.

Here we have used a query within a query—an ‘inner query’, or ‘subquery’.

```

1 PREFIX : <http://www.oxfordsemantic.tech/f1demo/>
2
3 # Most races run without ever getting a podium
4 SELECT ?forename ?surname ?raceCount
5 WHERE {
6 # First, get all the drivers with their names.
7 # The order of the query atoms is mostly not important, as RDFS will rearrange it
8 # internally in order to have the best performance automatically.
9   ?driver a :drivers ;
10      :drivers_forename ?forename ;
11      :drivers_surname ?surname .
12
13 # Then make sure that they have never achieved a podium (i.e. positions 3, 2, or 1)
14   FILTER NOT EXISTS {
15     ?result2 :results_driver ?driver ;
16     :results_positionOrder ?positionOrder .
17     FILTER(?positionOrder IN(1, 2, 3))
18   }
19
20 # Finally find out how many races they have raced in.
21 # This is done in an 'inner query'
22   {
23     SELECT ____ (COUNT(____) AS ?raceCount)
24     WHERE {
25       ?result :results_driver ?driver ;
26       :results_race ?race .
27     }
28     GROUP BY ____
29   }
30 }
31 ORDER BY DESC(____) # Don't forget to order

```



Query 11 – Driver Win Percentage

Calculate the **win percentage** of each driver individually, ranking them from most to least successful.

The **BIND** function sets a variable to a specific value.

```

1 PREFIX : <http://www.oxfordsemantic.tech/fldemo/>
2
3 # Drivers with their win percentage, ordered by win percentage
4 SELECT ?forename ?surname ?raceCount ?raceWins ?percentage
5 WHERE {
6   # First get the drivers
7   ?driver :drivers_forename ?forename ;
8           :drivers_surname ?surname .
9
10  # Then get the race count for each driver with an innery query...
11  {
12    SELECT ?driver (COUNT(?race) AS ?raceCount)
13    WHERE {
14      ?result :results_driver ?driver ;
15             :results_race ?race .
16    }
17    GROUP BY ?driver
18  }
19
20  # ... and get the *win* count for each driver with another inner query.
21  {
22    SELECT ___ (COUNT(?race) AS ___)
23    WHERE {
24      ?result :results_driver ___ ;
25             :results_race ___ ;
26             :results_positionOrder 1 .
27    }
28    GROUP BY ?driver
29  }
30
31  # Finally use the two aggregate variables to compute a percentage
32  # with the BIND keyword.
33  BIND(?raceWins/___ AS ?percentage)
34
35 }
36 ORDER BY DESC(?percentage)

```



Query 12 – Correcting Win Percentage

Correct the calculation of win percentages to include even those drivers who never won.

The **OPTIONAL** keyword states that if, for a given case, the subsequent result cannot be evaluated, skip over it and leave it undefined.

```

1 PREFIX : <http://www.oxfordsemantic.tech/f1demo/>
2
3 # Can you spot a problem?
4
5 # We only return 'biased' results, i.e. those who have at least one win!
6 # We will solve this using the so called OPTIONAL keyword
7
8 SELECT ?forename ?surname ?raceCount ?raceWinsFinal ?percentage
9 WHERE {
10 # First get the drivers
11 ?driver :drivers_forename ?forename ;
12         :drivers_surname ?surname .
13
14 # First get the race count for each driver.
15 {SELECT ?driver (COUNT(?race) AS ?raceCount)
16  WHERE {
17    ?result :results_driver ?driver ;
18            :results_race ?race .
19  }
20  GROUP BY ?driver}
21
22 # Then *optionally* get the win count for each driver.
23 OPTIONAL {SELECT ?driver (COUNT(?race) AS ?raceWins)
24  WHERE {
25    ?result :results_driver ?driver ;
26            :results_race ?race ;
27            :results_positionOrder 1 .
28  }
29  GROUP BY ?driver}
30
31 # Using the COALESCE keyword, we make sure that when the number of wins is undefined
32 # it is 'bound' as 0.
33 BIND(COALESCE(____,0) AS ?raceWinsFinal)
34 # And use the two aggregate variables to compute a percentage
35 BIND(____/?raceCount AS ?percentage)
36
37 }
38 ORDER BY DESC(____)

```

© Oxford Semantic Technologies 2022

The **COALESCE** function returns the first value in its list that does not throw up an error (including undefined).

It is bound to the stated variable.

Exercise 3



Complete **Queries 9-12**.

HINT:

To finish on the podium you must come in positions 1, 2, or 3.

Every result will be associated with a position number which tells us where they finished in the race.



Exercise 3 – Bonus Question

Use **FILTER** to restrict this list to only drivers who **won at least 5** races.

Answer



Running SPARQL from the command line

As well as running SPARQL queries from the web console, queries can be run from the command line.

Here are some of the commands that you will need to set this up:

`active <data store name>`

- The active command sets the active data store for subsequent commands.

`set output out`

- The set command allows you to change internal variables used by RDFox.
- `set output out` sets the output variable so that query results are sent to the terminal.

`evaluate <query file name>`

- Runs the query with that filename.



<https://docs.oxfordsemantic.tech/command-line-reference.html>



Running SPARQL from the command line

From the RDFox command line:

- a. Make the **f1** data store active.

```
active f1
```

- b. Set the output to the terminal.

```
set output out
```

- c. Run **query 6** again.

If you are running RDFox from **RDFoxWorkshop-SPARQL** directory then the command should be

```
evaluate Queries/q6.rq
```

Congratulations!



A look back at what you've achieved.

Setting up RDFox

- ✓ License and executable
- ✓ IDE
- ✓ REST endpoint and UI

Loading data into RDFox

- ✓ Create datastore
- ✓ Import data

Exploring and querying with SPARQL

- ✓ Basics of SPARQL
- ✓ Useful queries to explore
- ✓ Aggregates
- ✓ Negation
- ✓ Filters
- ✓ Binds
- ✓ Optionals

Closing RDFox



Stop your RDFox server with the `quit` command.



Your final exercise...

Click on the link below to start filling out our quick survey. It won't take you more than 3 minutes!

<https://oxfordsemantictech.typeform.com/to/webHJ4rs>

Further resources



Our website

<https://www.oxfordsemantic.tech>

Request an evaluation license

<https://www.oxfordsemantic.tech/tryrdfoxforfree>

Read the documentation

<https://docs.oxfordsemantic.tech/>

Our blog

<https://www.oxfordsemantic.tech/the-blog>

